

基于案例学SQL优化第2周

从案例中分析体系结构如何左右SQL性能

【声明】 本视频和幻灯片为炼数成金网络课程的教学资料，所有资料只能在课程内使用，不得在课程以外范围散播，违者将可能被追究法律和经济责任。

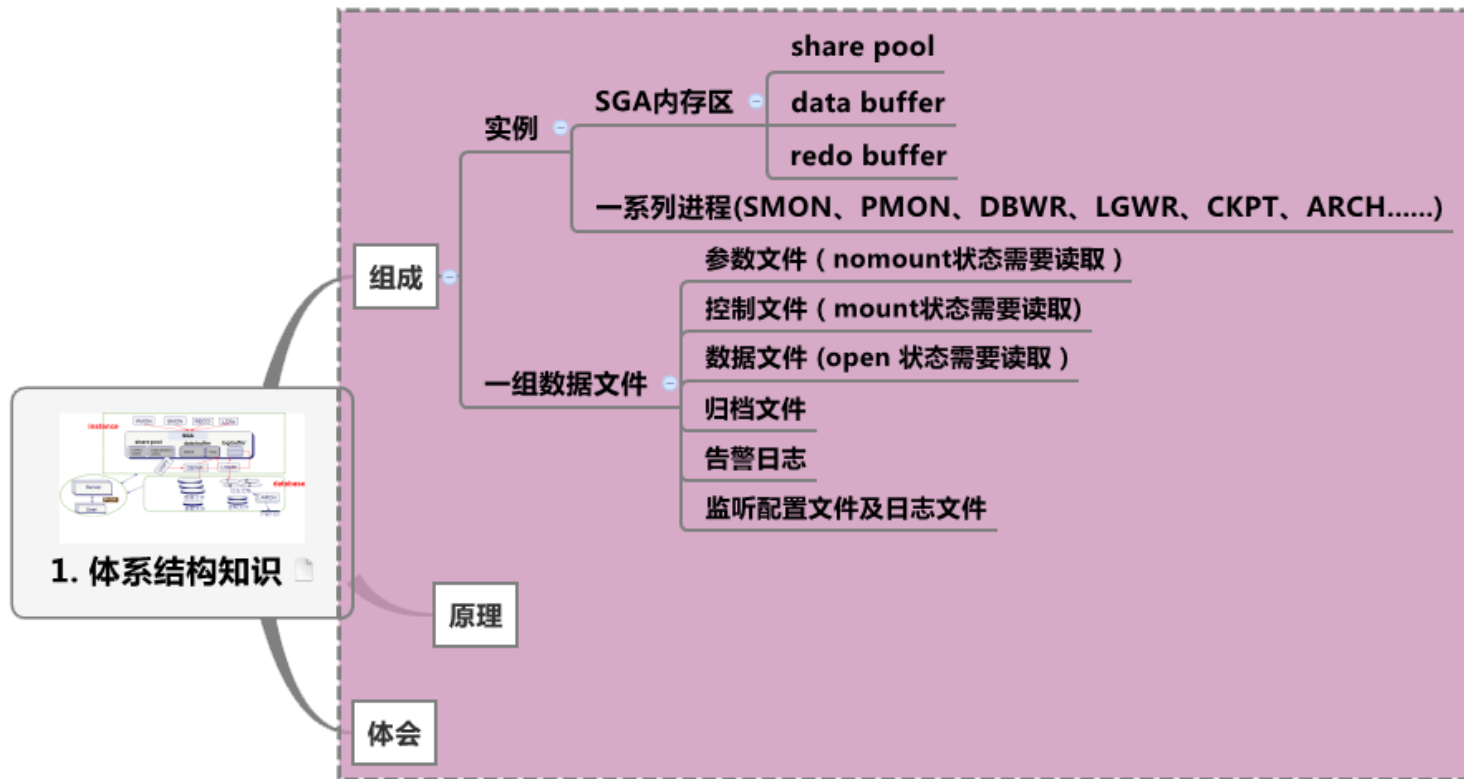
课程详情访问炼数成金培训网站

<http://edu.dataguru.cn>

基于案例学SQL优化

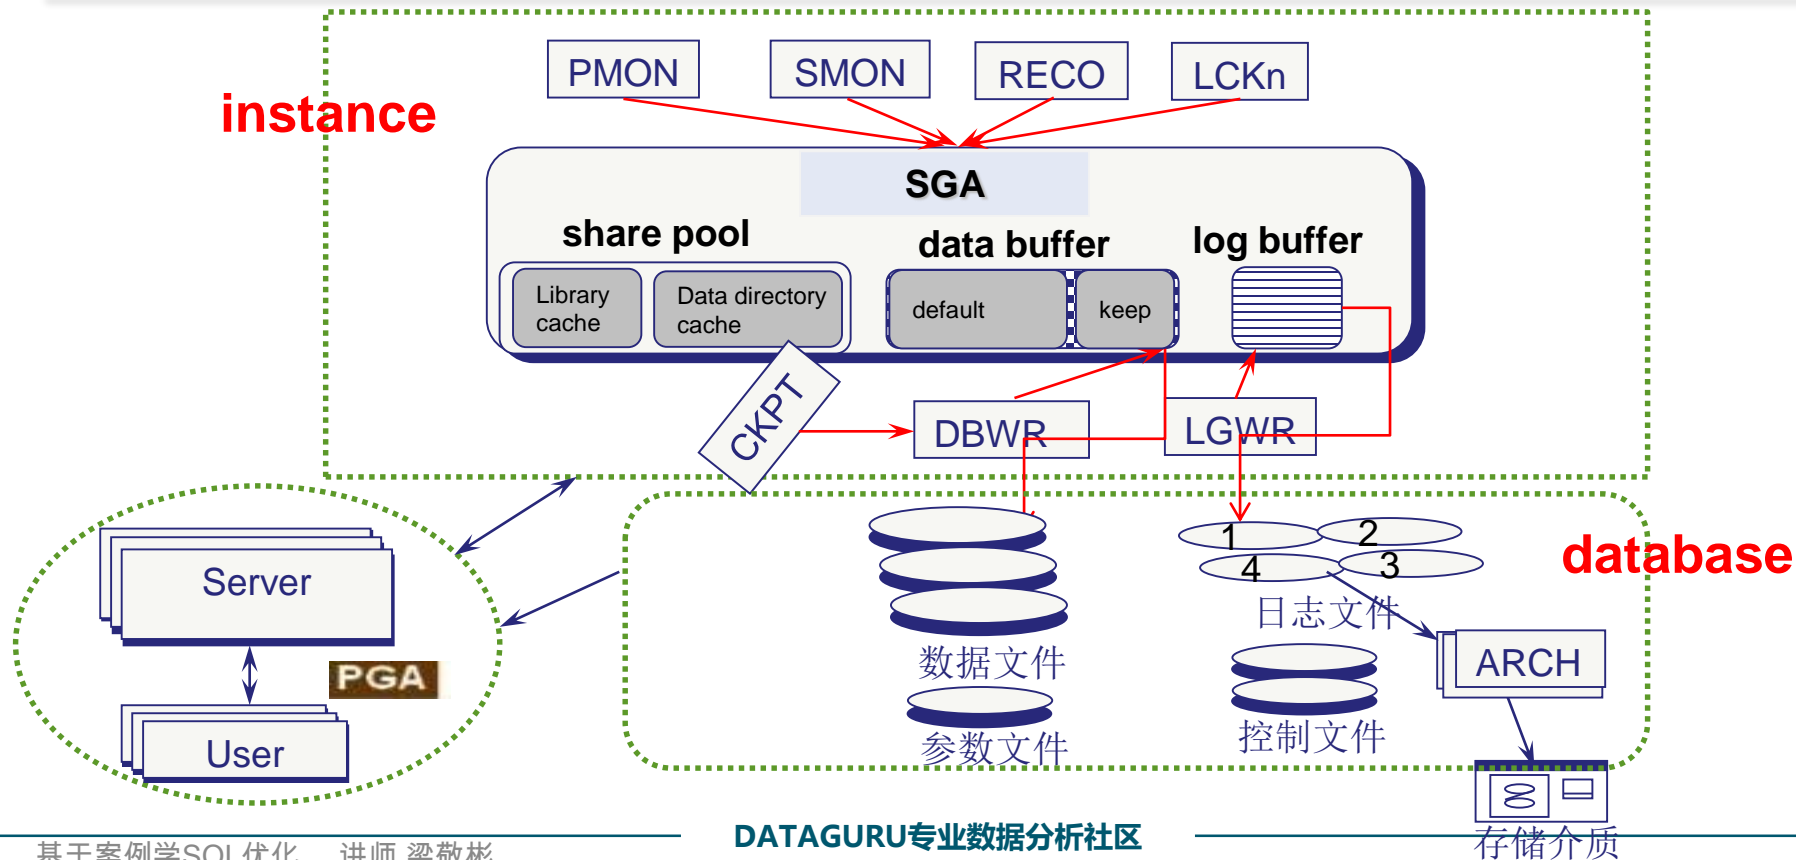
- 第1周 从案例中推导SQL优化的总体思路与误区
- 第2周 从案例中分析体系结构如何左右SQL性能
- 第3周 从案例中体验逻辑结构如何影响SQL优化
- 第4周 从案例中探寻表设计对SQL优化的重要性
- 第5周 从案例中明白索引是如何让SQL运行飞快
- 第6周 从案例中体会索引让SQL举步维艰的一面
- 第7周 从案例中体会函数及位图索引与SQL优化
- 第8周 从案例中洞察表连接与SQL优化之间关系
- 第9周 从案例中探讨该如何分析读懂析执行计划
- 第10周 从案例中学习左右Oracle执行计划之妙法
- 第11周 从案例中学会应用工具进行SQL整体优化
- 第12周 从案例中学习如何快速缩短SQL优化过程
- 第13周 从案例中感悟SQL等价改写优化思路之1
- 第14周 从案例中感悟SQL等价改写优化思路之2
- 第15周 从课程所有案例理出SQL优化思路及意识





Oracle的体系结构(简化版)

instance



1. 未启动数据库前的SGA分配情况

```
[oracle@itmapp3 ~]$ ipcs -m
----- Shared Memory Segments -----
key          shmid      owner      perms      bytes      nattch     status
0x00020014   32769      gserver    666        12         10
0x00028081   131076     kfv3       666        12         6
0x00024024   1015828    nmv3       666        12         6
```

2. 启动数据库后的SGA分配情况

```
[oracle@itmapp3 ~]$ ipcs -m
----- Shared Memory Segments -----
key          shmid      owner      perms      bytes      nattch
0x66a02988   1114112    oracle     660        2485125120 14
0x00020014   32769      gserver    666        12         10
0x00028081   131076     kfv3       666        12         6
0x00024024   1015828    nmv3       666        12         6
```

3. 为啥是2485125120字节，原来如此：

```
SQL> show parameter sga
NAME                                TYPE                                VALUE
-----
lock_sga                            boolean                             FALSE
pre_page_sga                        boolean                             FALSE
sga_max_size                         big integer                         2368M
sga_target                           big integer                         2368M
```

注：以上是10g的分配方式，在11g中默认是内存管理分配sga（将sga和pga一同管理）

```
SQL> show parameter memory
NAME                                TYPE                                VALUE
-----
hi_shared_memory_address            integer                             0
memory_max_target                    big integer                         4G
memory_target                        big integer                         4G
shared_memory_address                integer                             0
```

体会体系结构中的进程

1. 未启动数据库前的oracle进程情况

```
[oracle@itmapp3 ~]$ ps -ef |grep ora
oracle 6796 1 0 Oct28 ? 00:06:16 /home/oracle/product/10.2.0/db_1/bin/tnslsnr LISTENER -inherit
```

2. 启动数据库后的oracle本地进程

```
[oracle@itmapp3 ~]$ ps -ef |grep ora
oracle 5601 1 0 05:50 ? 00:00:00 ora_pmon_itmtest
oracle 5606 1 0 05:50 ? 00:00:00 ora_psp0_itmtest
oracle 5614 1 0 05:50 ? 00:00:00 ora_mman_itmtest
oracle 5616 1 0 05:50 ? 00:00:00 ora_dbw0_itmtest
oracle 5618 1 0 05:50 ? 00:00:00 ora_lgwr_itmtest
oracle 5620 1 0 05:50 ? 00:00:00 ora_ckpt_itmtest
oracle 5624 1 0 05:50 ? 00:00:00 ora_smon_itmtest
oracle 5626 1 0 05:50 ? 00:00:00 ora_reco_itmtest
oracle 5628 1 0 05:50 ? 00:00:00 ora_cjq0_itmtest
oracle 5630 1 0 05:50 ? 00:00:00 ora_mmon_itmtest
oracle 5632 1 0 05:50 ? 00:00:00 ora_mmln_itmtest
oracle 5634 1 0 05:50 ? 00:00:00 ora_d000_itmtest
oracle 5636 1 0 05:50 ? 00:00:00 ora_s000_itmtest
oracle 5637 5037 0 05:50 ? 00:00:00 oracleitmttest (DESCRIPTION=(LOCAL=YES) (ADDRESS=(PROTOCOL=beq)))
oracle 6796 1 0 Oct28 ? 00:06:17 /home/oracle/product/10.2.0/db_1/bin/tnslsnr LISTENER -inherit
```

3. 接下来用户连上来的进程情况

```
oracle 6799 1 0 05:34 ? 00:00:00 oracleitmttest (LOCAL=NO)
oracle 6807 1 0 05:34 ? 00:00:00 oracleitmttest (LOCAL=NO)
oracle 6811 1 0 05:34 ? 00:00:00 oracleitmttest (LOCAL=NO)
oracle 6820 1 0 05:34 ? 00:00:00 oracleitmttest (LOCAL=NO)
oracle 6822 1 0 05:34 ? 00:00:00 oracleitmttest (LOCAL=NO)
oracle 6824 1 0 05:34 ? 00:00:00 oracleitmttest (LOCAL=NO)
oracle 8801 1 0 Nov04 ? 00:00:00 oracleitmttest (LOCAL=NO)
oracle 8803 1 0 Nov04 ? 00:00:00 oracleitmttest (LOCAL=NO)
```

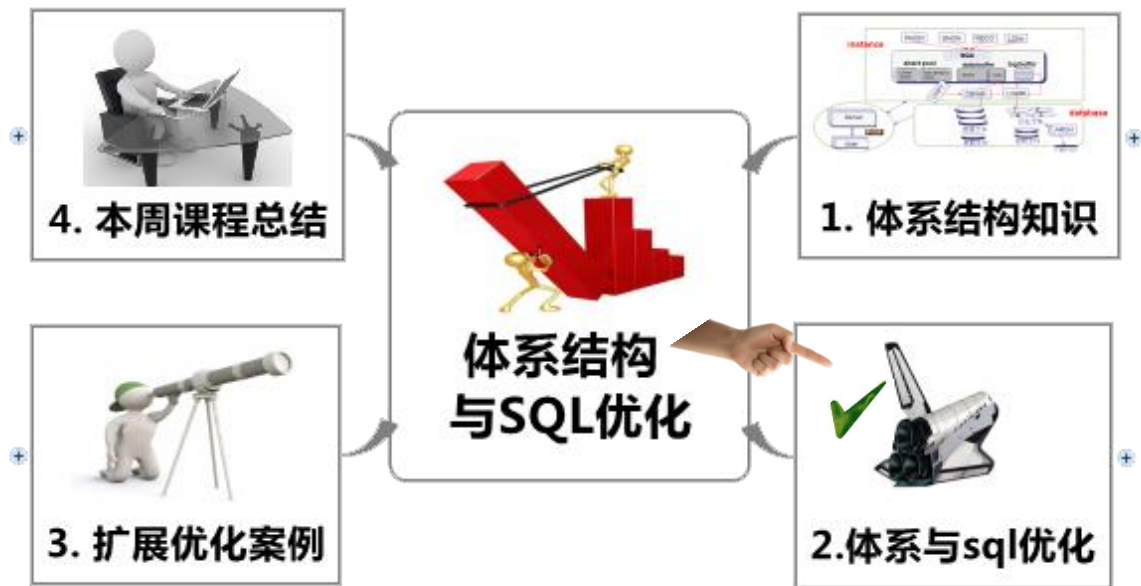
4. Oracle本地进程和外部连接进程统计

```
[oracle@itmapp3 ~]$ ps -ef |grep LOCAL=NO|wc -l
```

150

```
[oracle@itmapp3 ~]$ ps -ef |grep ora|wc -l
```

172





2.体系与sql优化 M

与共享池相关

解析优化让第2次执行更快 📄

思考绑定变量带来性能飞跃 📄

体会绑定变量的AWR试验 📄

体会参数对SQL性能的影响 📄

数据缓冲相关

缓冲优化让第2次执行更快 📄

直接路径读性能略胜一筹 📄

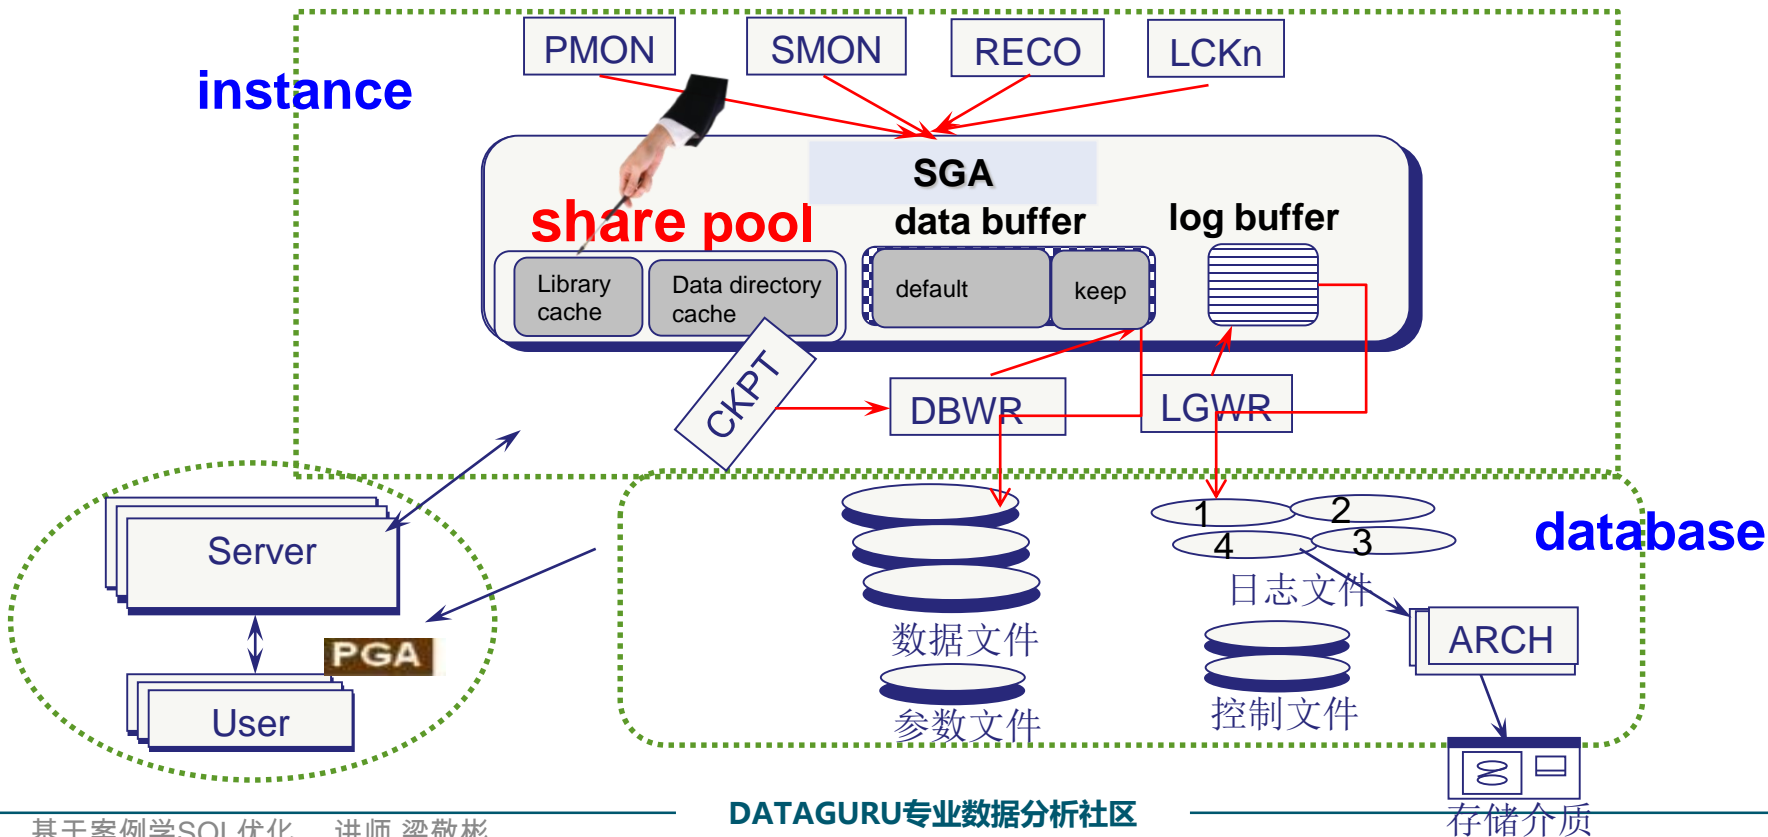
日志归档相关

批量提交与否性能差异明显 📄

日志关闭与否对性能的影响 📄

体系结构中共享池研究

instance



解析优化让第2次执行更快

SQL> --第1次执行

```
SQL> select count(*) from t;  
COUNT(*)
```

72884

已用时间: 00: 00: 00.10

执行计划

Id	Operation	Name	Rows	Cost (%CPU)	Time
0	SELECT STATEMENT		1	291 (1)	00:00:04
1	SORT AGGREGATE		1		
2	TABLE ACCESS FULL	T	60918	291 (1)	00:00:04

统计信息

```
28 recursive calls  
0 db block gets  
1103 consistent gets  
1038 physical reads  
0 redo size  
425 bytes sent via SQL*Net to client  
415 bytes received via SQL*Net from client  
2 SQL*Net roundtrips to/from client  
0 sorts (memory)  
0 sorts (disk)  
1 rows processed
```

SQL> --第2次执行

SQL> --该命令只是为了先不考虑2次执行物理读减少带来的效果，只考虑减少解析的优化效果
SQL> alter system flush buffer_cache;
系统已更改。

已用时间: 00: 00: 00.03

```
SQL> select count(*) from t;  
COUNT(*)
```

72884

已用时间: 00: 00: 00.07

执行计划

Id	Operation	Name	Rows	Cost (%CPU)	Time
0	SELECT STATEMENT		1	291 (1)	00:00:04
1	SORT AGGREGATE		1		
2	TABLE ACCESS FULL	T	60918	291 (1)	00:00:04

统计信息

```
0 recursive calls  
0 db block gets  
1043 consistent gets  
1039 physical reads  
0 redo size  
425 bytes sent via SQL*Net to client  
415 bytes received via SQL*Net from client  
2 SQL*Net roundtrips to/from client  
0 sorts (memory)  
0 sorts (disk)
```

思考绑定变量所带来的性能飞跃

```
SQL> begin
  2   for i in 1 .. 100000
  3   loop
  4       execute immediate
  5       'insert into t values ( '||i||' )';
  6   end loop;
  7   commit;
  8 end;
  9 /
```

PL/SQL 过程已成功完成。

已用时间: 00: 00: 43.50

SQL>--使用绑定变量

```
SQL> begin
  2   for i in 1 .. 100000
  3   loop
  4       execute immediate
  5       'insert into t values ( :x )' using i;
  6   end loop;
  7   commit;
  8 end;
  9 /
```

PL/SQL 过程已成功完成。

已用时间: 00: 00: 04.77

共享池相关的优化


自行体会硬解析次数和执行次数

SQL

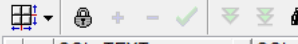
Output

Statistics

select t.sql_text, t.sql_id, t.executions, t.parse_calls
from v\$sql t
where sql_text like 'insert into t values%';



	SQL_TEXT	SQL_ID	EXECUTIONS	PARSE_CALLS
1	insert into t values (96283) ...	b5qk5k0cg000f	1	1
2	insert into t values (95426) ...	ch0act8w48019	1	1
3	insert into t values (94549) ...	0v81tsyp1401h	1	1
4	insert into t values (96599) ...	2gn81fnjd001m	1	1
5	insert into t values (94842) ...	7h9y8au96n02g	1	1
6	insert into t values (96579) ...	b8ubtwdrac02p	1	1
7	insert into t values (92824) ...	aw6f4nsnt0032	1	1
8	insert into t values (92398) ...	1ds5nr9nfw034	1	1
9	insert into t values (95685) ...	1uuu6w9ykn034	1	1
10	insert into t values (96192) ...	g3cfj18vh8039	1	1
11	insert into t values (98214) ...	bsv2kq808403m	1	1
12	insert into t values (96354) ...	4zrq88kw003w	1	1
13	insert into t values (98195) ...	gmu01syv2c04y	1	1
14	insert into t values (98652) ...	agvzvxxvm4056	1	1
15	insert into t values (97683) ...	0nvm3ta8t005f	1	1
16	insert into t values (96015) ...	147a1adk4a061	1	1

SQL	Output	Statistics		
<pre>select t.sql_text, t.sql_id, t.executions, t.parse_calls from v\$sql t where sql_text like 'insert into t values (:x)%';</pre>				
				
	SQL_TEXT	SQL_ID	EXECUTIONS	PARSE_CALLS
▶ 1	insert into t values (x) ...	bbz6pdq577unj	100000	1

共享池相关的优化

体会绑定变量的AWR试验

	Per Second	Per Transaction	Per Exec	Per Call
DB Time(s):	0.9	3.2	0.00	3.17
DB CPU(s):	0.8	3.1	0.00	3.14
Redo size:	436,015.3	1,631,772.8		
Logical reads:	7,943.7	29,729.2		
Block changes:	3,625.5	13,568.4		
Physical reads:	8.0	30.0		
Physical writes:	0.3	1.2		
User calls:	0.3	1.0		
Parses:	1,823.4	6,824.0		
Hard parses:	1,790.5	6,701.0		
W/A MB processed:	0.5	1.9		
Logons:	0.1	0.2		
Executes:	1,931.9	7,229.9		
Rollbacks:	0.0	0.0		
Transactions:	0.3			

Buffer Nowait %:	100.00	Redo NoWait %:	100.00
Buffer Hit %:	99.90	In-memory Sort %:	100.00
Library Hit %:	62.91	Soft Parse %:	1.80
Execute to Parse %:	5.61	Latch Hit %:	100.00
Parse CPU to Parse Elapsed %:	98.37	% Non-Parse CPU:	25.58

体会绑定变量的trace试验

--未使用绑定变量的

OVERALL TOTALS FOR ALL RECURSIVE STATEMENTS

call	count	cpu	elapsed	disk	query	current	rows
Parse	10075	3.32	3.39	0	0	0	0
Execute	10379	0.60	0.56	11	10049	30380	10006
Fetch	683	0.01	0.10	219	1565	0	1418
total	21137	3.94	4.06	230	11614	30380	11424

---使用绑定变量的

OVERALL TOTALS FOR ALL RECURSIVE STATEMENTS

call	count	cpu	elapsed	disk	query	current	rows
Parse	76	0.03	0.02	0	0	0	0
Execute	10379	1.41	1.45	12	52	10397	10006
Fetch	683	0.06	0.09	219	1565	0	1418
total	11138	1.51	1.58	231	1617	10397	11424

注意静态SQL自动绑定变量

```
SQL> set linesize 1000
SQL> column sql_text format a50
SQL> select t.sql_text, t.sql_id, t.executions, t.parse_calls
       2   from v$sql t
       3   where lower(sql_text) like 'insert into t values%';
```

SQL_TEXT	SQL_ID	EXECUTIONS	PARSE_CALLS
-----	-----	-----	-----
INSERT INTO T VALUES (:B1)	2n9c7yutww4dx4	100000	0

体会参数对SQL性能的影响

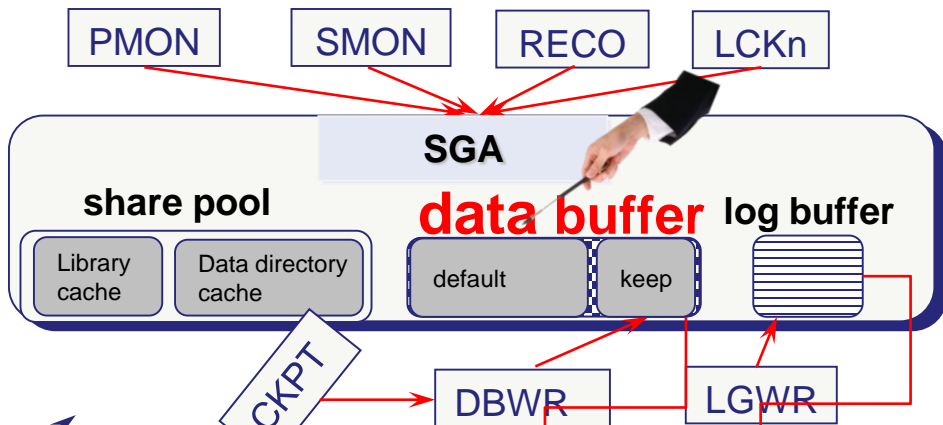


```
SQL> set timing on
SQL> alter session set session_cached_cursors=0;
会话已更改。
SQL> --使用绑定变量
SQL> begin
  2     for i in 1 .. 100000
  3     loop
  4         execute immediate
  5             'insert into t values ( :x )' using i;
  6     end loop;
  7     commit;
  8 end;
  9 /
PL/SQL 过程已成功完成。
已用时间: 00: 00: 08.70
```

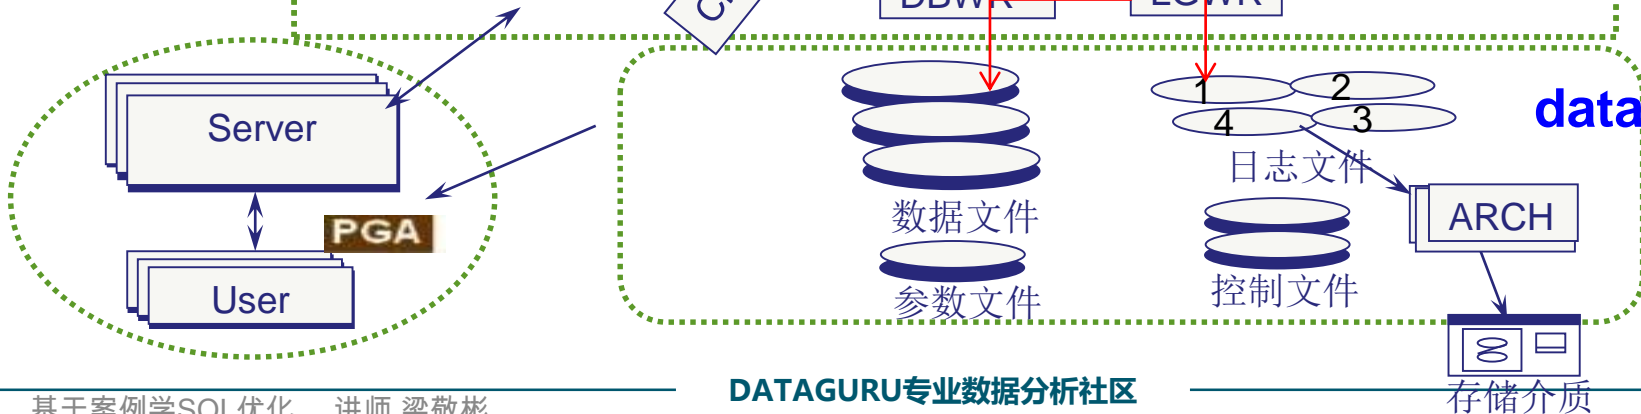
```
SQL> set timing on
SQL> alter session set session_cached_cursors=50;
SQL> --使用绑定变量
SQL> begin
  2     for i in 1 .. 100000
  3     loop
  4         execute immediate
  5             'insert into t values ( :x )' using i;
  6     end loop;
  7     commit;
  8 end;
  9 /
PL/SQL 过程已成功完成。
已用时间: 00: 00: 04.69
```

体系结构中的数据缓冲池研究

instance



database





缓冲优化让第2次执行更快

```
SQL> --第1次执行
SQL> select count(*) from t;
COUNT(*)
```

72884

已用时间: 00: 00: 00.12

执行计划

Id	Operation	Name	Rows	Cost (%CPU)	Time
0	SELECT STATEMENT		1	291 (1)	00:00:04
1	SORT AGGREGATE		1		
2	TABLE ACCESS FULL	T	60918	291 (1)	00:00:04

统计信息

```
28 recursive calls
0 db block gets
1103 consistent gets
1038 physical reads
0 redo size
425 bytes sent via SQL*Net to client
415 bytes received via SQL*Net from client
2 SQL*Net roundtrips to/from client
0 sorts (memory)
0 sorts (disk)
1 rows processed
```

```
SQL> --第2次执行
SQL> --该命令只是为了先不考虑解析的优化,单纯考虑第2次执行物理读减少带来的优化效应
SQL> alter system flush shared_pool;
系统已更改.
```

```
SQL> select count(*) from t;
COUNT(*)
```

72884

已用时间: 00: 00: 00.04

执行计划

Id	Operation	Name	Rows	Cost (%CPU)	Time
0	SELECT STATEMENT		1	291 (1)	00:00:04
1	SORT AGGREGATE		1		
2	TABLE ACCESS FULL	T	60918	291 (1)	00:00:04

统计信息

```
282 recursive calls
0 db block gets
1131 consistent gets
0 physical reads
0 redo size
425 bytes sent via SQL*Net to client
415 bytes received via SQL*Net from client
2 SQL*Net roundtrips to/from client
5 sorts (memory)
0 sorts (disk)
1 rows processed
```

解析和缓冲优化正常一起来

```
SQL> --第1次执行
SQL> select count(*) from t;
COUNT(*)
```

72884

已用时间: 00: 00: 00.11

执行计划

Id	Operation	Name	Rows	Cost (%CPU)	Time
0	SELECT STATEMENT		1	291 (1)	00:00:04
1	SORT AGGREGATE		1		
2	TABLE ACCESS FULL	T	69684	291 (1)	00:00:04

统计信息

```
28 recursive calls
0 db block gets
1110 consistent gets
1038 physical reads
0 redo size
425 bytes sent via SQL*Net to client
415 bytes received via SQL*Net from client
2 SQL*Net roundtrips to/from client
0 sorts (memory)
0 sorts (disk)
1 rows processed
```

```
SQL> --第2次执行
SQL> --这里不做 shared_pool和buffer_cache的flush
SQL> select count(*) from t;
COUNT(*)
```

72884

已用时间: 00: 00: 00.02

执行计划

Id	Operation	Name	Rows	Cost (%CPU)	Time
0	SELECT STATEMENT		1	291 (1)	00:00:04
1	SORT AGGREGATE		1		
2	TABLE ACCESS FULL	T	69684	291 (1)	00:00:04

统计信息

```
0 recursive calls
0 db block gets
1043 consistent gets
0 physical reads
0 redo size
425 bytes sent via SQL*Net to client
415 bytes received via SQL*Net from client
2 SQL*Net roundtrips to/from client
0 sorts (memory)
0 sorts (disk)
1 rows processed
```

数据缓冲相关优化

直接路径读性能略胜一筹

```
SQL> --测试普通插入
SQL> set timing on
SQL> insert into test select * from t;
已创建1166096行。
已用时间: 00: 00: 06.78
```

```
SQL> --测试直接路径读方式
SQL> set timing on
SQL> insert /* append */ into test select * from t;
已创建1166096行。
已用时间: 00: 00: 01.24
```

能胜一筹是有原因的

```
SQL> --注意普通方式插入test表后输出的物理读（首次执行，看physical reads）
SQL> set autotrace traceonly
SQL> select count(*) from test;
```

Id	Operation	Name	Rows	Cost (%CPU)	Time
0	SELECT STATEMENT		1	4685 (1)	00:00:57
1	SORT AGGREGATE		1		
2	TABLE ACCESS FULL	TEST	1255K	4685 (1)	00:00:57

统计信息

```

29 recursive calls
1 db block gets
17008 consistent gets
0 physical reads
176 redo size
426 bytes sent via SQL*Net to client
415 bytes received via SQL*Net from client
2 SQL*Net roundtrips to/from client
0 sorts (memory)
0 sorts (disk)
1 rows processed
```

```
SQL> --注意直接路径方式插入test表试验输出的物理读（首次执行，看physical reads）
SQL> set autotrace traceonly
SQL> select count(*) from test;
```

执行计划

Id	Operation	Name	Rows	Cost (%CPU)	Time
0	SELECT STATEMENT		1	4559 (1)	00:00:55
1	SORT AGGREGATE		1		
2	TABLE ACCESS FULL	TEST	1052K	4559 (1)	00:00:55

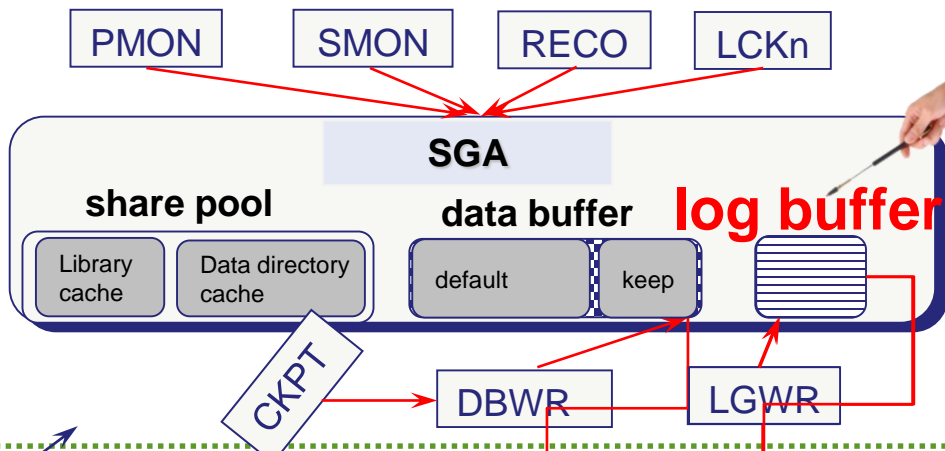
统计信息

```

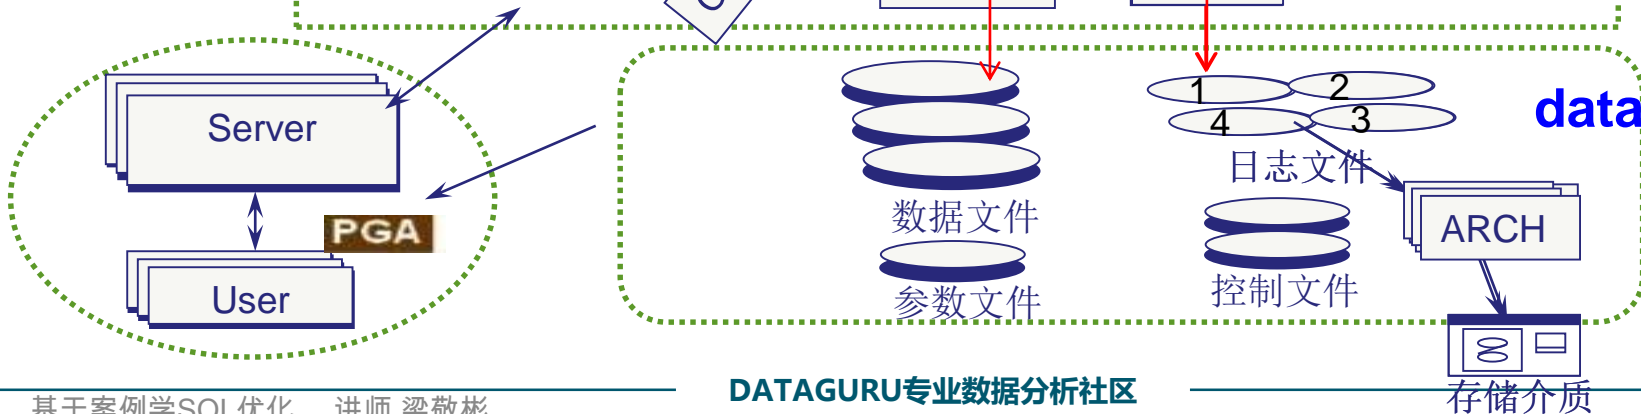
28 recursive calls
1 db block gets
16688 consistent gets
16604 physical reads
168 redo size
426 bytes sent via SQL*Net to client
415 bytes received via SQL*Net from client
2 SQL*Net roundtrips to/from client
0 sorts (memory)
0 sorts (disk)
1 rows processed
```

体系结构中的日志相关研究

instance



database



批量提交与否性能差异明显



```
SQL> drop table t purge;
表已删除。
SQL> create table t(x int);
表已创建。
SQL> set timing on
SQL> begin
2       for i in 1 .. 100000 loop
3           insert into t1 values (i);
4           commit;
5       end loop;
6 end;
7 /
PL/SQL 过程已成功完成。
已用时间: 00: 00: 11.21
```

```
SQL> drop table t purge;
表已删除。
SQL> create table t(x int);
表已创建。
SQL> begin
2       for i in 1 .. 100000 loop
3           insert into t values (i);
4       end loop;
5       commit;
6 end;
7 /
PL/SQL 过程已成功完成。
已用时间: 00: 00: 04.26
```

日志关闭与否对性能的影响

```
SQL> --测试直接路径读方式
SQL> drop table test;
表已删除。
SQL> create table test as select * from dba_objects where 1=2;
表已创建。
SQL> set timing on
SQL> insert /*+ append */ into test select * from t;
已创建4664384行。
已用时间: 00: 00: 05.01
```

```
SQL> --测试nologging关闭日志+直接路径读方式
SQL> drop table test;
表已删除。
SQL> create table test as select * from dba_objects where 1=2;
表已创建。
SQL> alter table test nologging;
表已更改。
SQL> set timing on
SQL> insert /*+ append */ into test select * from t;
已创建4664384行。
已用时间: 00: 00: 04.39
```





3. 扩展优化案例

与共享池相关

头疼，如何查硬解析问题

纠结，绑定变量也有冬天

经典，SQL的逻辑读变零

经典，函数的逻辑读成零

数据缓冲相关

感谢，keep让sql跑更快

细致，查系统各维度规律

日志归档相关

规律，日志切换有据可查

巧妙，逮到提交过频语句

头疼，如何查硬解析问题



```
SQL> ---执行完上述动作后，以下SQL语句可以完成未绑定变量语句的统计
SQL> set linesize 266
SQL> col sql_text_wo_constants format a30
SQL> col module format a30
SQL> col CNT format 999999
SQL> select sql_text_wo_constants, module, count(*) CNT
  2   from t_bind_sql
  3   group by sql_text_wo_constants, module
  4   having count(*) > 100
  5   order by 3 desc;
```

SQL_TEXT_WO_CONSTANTS	MODULE	CNT
INSERT INTO T VALUES (@)	SQL*Plus	7366

共享池的相关案例

纠结，绑定变量也有冬天

SQL> SELECT count(pad) FROM t WHERE id < 990;

0	SELECT STATEMENT		1	105	7	(0)	00:00:01
1	SORT AGGREGATE		1	105			
* 2	TABLE ACCESS FULL	T	990	101K	7	(0)	00:00:01

SQL> SELECT count(pad) FROM t WHERE id < 10;

0	SELECT STATEMENT		1	105	3	(0)	00:00:01
1	SORT AGGREGATE		1	105			
2	TABLE ACCESS BY INDEX ROWID	T	9	945	3	(0)	00:00:01
* 3	INDEX RANGE SCAN	T_PK	9		2	(0)	00:00:01

SQL> EXECUTE :id := 990;

PL/SQL 过程已成功完成。

SQL> SELECT count(pad) FROM t WHERE id < :id;

0	SELECT STATEMENT		1	105	3	(0)	00:00:01
1	SORT AGGREGATE		1	105			
2	TABLE ACCESS BY INDEX ROWID	T	50	5250	3	(0)	00:00:01
* 3	INDEX RANGE SCAN	T_PK	9		2	(0)	00:00:01

SQL> EXECUTE :id := 10;

SQL> SELECT count(pad) FROM t WHERE id < :id;

0	SELECT STATEMENT		1	105	3	(0)	00:00:01
1	SORT AGGREGATE		1	105			
2	TABLE ACCESS BY INDEX ROWID	T	50	5250	3	(0)	00:00:01
* 3	INDEX RANGE SCAN	T_PK	9		2	(0)	00:00:01

共享池的相关案例

经典，SQL的逻辑读变零

```
SQL> ---接下来再次执行(居然发现逻辑读为0):
SQL> set autotrace on
SQL> select /*+ result_cache */ count(*) from t;
COUNT(*)
-----
145762
```

已用时间: 00:00:00.01

执行计划

Id	Operation	Name	Rows	Cost (%CPU)	Time
0	SELECT STATEMENT		1	589 (1)	00:00:08
1	RESULT CACHE	d827qx1jmwjc86yqynrplkvpy			
2	SORT AGGREGATE		1		
3	TABLE ACCESS FULL	T	277K	589 (1)	00:00:08

统计信息

```
0 recursive calls
0 db block gets
0 consistent gets
0 physical reads
0 redo size
425 bytes sent via SQL*Net to client
416 bytes received via SQL*Net from client
2 SQL*Net roundtrips to/from client
0 sorts (memory)
0 sorts (disk)
1 rows processed
```



经典，函数的逻辑读成零

```
SQL> --看调用F_NO_RESULT_CACHE执行第2次后的结果
SQL> SELECT F_NO_RESULT_CACHE FROM DUAL;
```

F_NO_RESULT_CACHE

72883

统计信息

1	recursive calls
0	db block gets
1043	consistent gets
0	physical reads
0	redo size
434	bytes sent via SQL*Net to client
415	bytes received via SQL*Net from client
2	SQL*Net roundtrips to/from client
0	sorts (memory)
0	sorts (disk)
1	rows processed

```
SQL> --看调用F_RESULT_CACHE执行第2次后的结果
SQL> SELECT F_RESULT_CACHE FROM DUAL;
```

F_RESULT_CACHE

72883

统计信息

0	recursive calls
0	db block gets
0	consistent gets
0	physical reads
0	redo size
431	bytes sent via SQL*Net to client
415	bytes received via SQL*Net from client
2	SQL*Net roundtrips to/from client
0	sorts (memory)
0	sorts (disk)
1	rows processed

数据缓冲相关案例

感谢，keep让sql跑更快

--未执行KEEP命令，通过如下查询出BUFFER_POOL列值为DEFAULT，表示未KEEP。

```
select BUFFER_POOL from user_tables where TABLE_NAME='T';
BUFFER_POOL
```

DEFAULT

```
select BUFFER_POOL from user_indexes where INDEX_NAME='IDX_OBJECT_ID';
```

```
BUFFER_POOL
```

DEFAULT

--执行KEEP操作后，通过如下查询出BUFFER_POOL列值为KEEP,表示已经KEEP成功了

```
select BUFFER_POOL from user_tables where TABLE_NAME='T';
```

```
BUFFER_POOL
```

KEEP

```
select BUFFER_POOL from user_indexes where INDEX_NAME='IDX_OBJECT_ID';
```

```
BUFFER_POOL
```

KEEP

数据缓冲相关案例

细致，查系统各维度规律

	SNAP_DATE	TIME	elapsed(min)	DB time(min)	REDO	redo/s	LOGICAL	logical/s	PHYSICAL	phy/s	EXECS	execs/s	PARSE	parse/s	HARDF
1	13/11/07	18:00	64.22	...	0.51	105873536	27478.21	1375950	357.11	225	0.06	160181	41.57	9018	2.34
2	13/11/07	16:56	115.77	...	0.04	1622588	233.6	43732	6.3	11	0	7177	1.03	4045	0.58
3	13/11/07	15:00	60.18	...	0.06	5690688	1575.93	106134	29.39	5518	1.53	13346	3.7	7542	2.09
4	13/11/07	14:00	60.17	...	0.65	197937568	54830.35	2338751	647.85	7149	1.98	406318	112.55	10217	2.83
5	13/11/07	13:00	60.18	...	0.21	126920828	35148.39	666700	184.63	6791	1.88	14344	3.97	9617	2.66
6	13/11/07	12:00	59.17	...	0.9	357269612	100639.33	1655494	466.34	6976	1.97	15255	4.3	9805	2.76
7	13/11/07	11:00	60.17	...	0.57	60423804	16737.9	310342	85.97	4853	1.34	553894	153.43	9828	2.72
8	13/11/07	08:01	60.18	...	0.06	2089688	578.7	59616	16.51	136	0.04	10912	3.02	6564	1.82
9	13/11/07	07:00	60.17	...	0.05	2543732	704.63	68438	18.96	469	0.13	12292	3.4	7171	1.99
10	13/11/07	06:00	34.1	...	0.62	3908728	1910.42	2313056	1130.53	2890	1.41	41855	20.46	8246	4.03
11	13/11/08	22:00	60.18	...	1.25	4584236	1269.52	500443	138.59	9367	2.59	50851	14.08	10969	3.04
12	13/11/08	21:00	60.17	...	0.06	2188264	606.17	57236	15.85	10	0	10882	3.01	6581	1.82
13	13/11/08	20:00	59.18	...	0.06	2340524	659.12	59905	16.87	11	0	10844	3.05	6570	1.85
14	13/11/08	19:01	60.9	...	0.09	2363340	646.78	55629	15.22	26	0.01	9964	2.73	5928	1.62
15	13/11/08	18:00	59.18	...	0.06	2395280	674.54	64119	18.06	40	0.01	11448	3.22	6930	1.95

巧妙，逮到提交过频语句

```
SQL> select t1.sid, t1.value, t2.name
       2   from v$sesstat t1, v$statname t2
       3   where t2.name like '%user commits%'
       4     and t1.STATISTIC# = t2.STATISTIC#
       5     and value >= 10000
       6   order by value desc;
SID          VALUE NAME
```

```
-----
194          100000 user commits
```

```
SQL> select t.SID, t.PROGRAM, t.EVENT, t.LOGON_TIME, t.WAIT_TIME, t.SECONDS_IN_WAIT, t.SQL_ID, t.PREV_SQL_ID
       from v$session t
       where sid in(194) ;
```

SID	PROGRAM	EVENT	LOGON_TIME	WAIT_TIME	SECONDS_IN_WAIT	SQL_ID	PREV_SQL_ID
194	sqlplus.exe	SQL*Net message from client	13-11月-13	0	54	ccpn5c32bmfmf	

```
SQL> select t.sql_id, t.sql_text, t.EXECUTIONS, t.FIRST_LOAD_TIME, t.LAST_LOAD_TIME
       from v$sql t
       where sql_id in ('ccpn5c32bmfmf');
```

SQL_ID	SQL_TEXT	EXECUTIONS	FIRST_LOAD_TIME	LAST_LOAD_TIME
ccpn5c32bmfmf	begin for i in 1 .. 100000 loop insert into t values (i); commit; end loop; end;	1	2013-11-13/14:41:18	2013-11-13/14:46:34

日志归档相关案例

规律，日志切换有据可查



	DAY	H00	H01	H02	H03	H04	H05	H06	H07	H08	H09	H10	H11	H12	H13	H14	H15	H16	H17	H18	H19	H20	H21	H22	H23	TOTAL
1	11/13	0	0	0	0	0	0	0	1	0	1	0	0	0	3	9	2	25	18	1	0	5	0	0	0	65
2	11/12	0	0	0	0	1	4	1	0	7	8	0	1	17	36	0	0	1	5	1	0	1	0	1	0	84
3	11/11	0	0	0	0	0	0	0	1	0	1	0	0	0	0	1	0	0	0	1	0	0	0	1	0	5
4	11/10	0	0	0	0	0	0	1	0	0	0	8	0	0	0	0	0	0	1	1	0	0	0	0	0	11
5	11/09	0	0	0	0	0	0	1	0	0	0	0	0	1	0	0	1	1	0	1	0	1	0	0	0	6
6	11/08	1	0	0	0	0	0	1	0	0	2	3	0	0	1	1	0	0	0	0	0	0	0	2	0	11
7	11/07	0	0	0	0	0	1	0	0	5	0	2	8	3	5	0	0	0	3	1	0	1	0	0	1	30
8	11/06	0	0	0	0	0	1	0	0	1	0	0	0	1	0	0	0	1	0	4	0	1	0	2	0	11
9	11/05	0	0	0	0	0	0	1	0	1	0	0	0	0	0	0	0	0	1	0	0	0	0	1	0	4
10	11/04	0	0	0	0	0	0	1	0	1	0	0	0	0	0	0	0	0	0	1	0	0	0	1	0	4

日志归档相关案例

迷案，跟踪日志暴增故障

---执行了大量针对test_redo表的INSERT操作后，，我们开始根据如下方法进行跟踪，看能否找到更新语句及对应的表。

--执行了大量的针对test_redo表的INSERT操作后，，我们开始根据如下方法进行跟踪，看能否发现更新的是哪张表，是哪些语句。

```
SQL> select * from (
  2  SELECT to_char(begin_interval_time, 'YYYY_MM_DD HH24:MI') snap_time,dhssso.object_name,SUM(db_block_changes_delta)
  3    FROM dba_hist_seg_stat dhss,dba_hist_seg_stat_obj dhssso,dba_hist_snapshot dhs
  4   WHERE dhs.snap_id = dhss. snap_id
  5         AND dhs.instance_number = dhss. instance_number AND dhss.obj# = dhssso. obj# AND dhss.dataobj# = dhssso.dataobj#
  6         AND begin_interval_time> sysdate - 60/1440
  7   GROUP BY to_char(begin_interval_time, 'YYYY_MM_DD HH24:MI'), dhssso.object_name order by 3 desc)
  8  where rownum<=3;
```

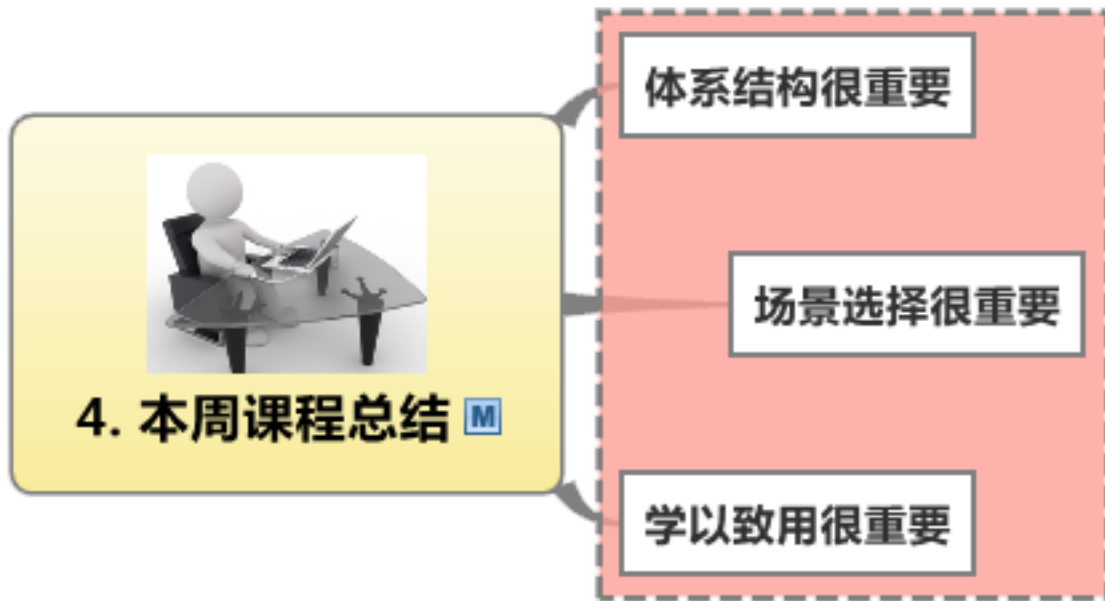
SNAP_TIME	OBJECT_NAME	SUM(DB_BLOCK_CHANGES_DELTA)
2013_11_13 20:00	TEST REDO	178272
2013_11_13 20:00	MGMT_CURRENT_METRICS_PK	224
2013_11_13 20:00	MGMT_SYSTEM_PERF_LOG_IDX_01	160

```
SQL> SELECT to_char(begin_interval_time,'YYYY_MM_DD HH24:MI'),dbms_lob.substr(sql_text,4000,1),dhss.sql_id,executions_delta,rows_processed_delta
  2  FROM dba_hist_sqlstat dhss, dba_hist_snapshot dhs, dba_hist_sqltext dhst
  3   WHERE UPPER(dhst.sql_text) LIKE '%TEST REDO%' AND dhss.snap_id = dhs.snap_id
  4   AND dhss.instance_number = dhss.instance_number AND dhss.sql_id = dhst.sql_id;
TO_CHAR(BEGIN_IN  DBMS_LOB.SUBSTR(SQL_TEXT,4000,1)          INSTANCE_NUMBER SQL_ID  EXECUTIONS_DELTA  ROWS_PROCESSED_DELTA
```

TO_CHAR(BEGIN_IN	DBMS_LOB.SUBSTR(SQL_TEXT,4000,1)	INSTANCE_NUMBER	SQL_ID	EXECUTIONS_DELTA	ROWS_PROCESSED_DELTA
2013 11 13 17:00	create table test redo as select * from dba objects	1	dsf2uj3pzzadg	1	72884
2013 11 13 20:00	insert into test redo select * from test redo	1	5w8pb7t27c85n	5	2259404

本周课程总结





炼数成金逆向收费式网络课程

- Dataguru (炼数成金) 是专业数据分析网站，提供教育，媒体，内容，社区，出版，数据分析业务等服务。我们的课程采用新兴的互联网教育形式，独创地发展了逆向收费式网络培训课程模式。既继承传统教育重学习氛围，重竞争压力的特点，同时又发挥互联网的威力打破时空限制，把天南地北志同道合的朋友组织在一起交流学习，使到原先孤立的学习个体组合成有组织的探索力量。并且把原先动辄成千上万的学习成本，直线下降至百元范围，造福大众。我们的目标是：低成本传播高价值知识，构架中国第一的网上知识流转阵地。
- 关于逆向收费式网络的详情，请看我们的培训网站 <http://edu.dataguru.cn>

Thanks

FAQ时间